This below is exactly what the app is about, end to end. Make sure the LLM models you are using are on the same page as what is below. The tech stack is also at the very bottom of the document so make sure you don't deviate.

What Remember is

A private web app where someone who has lost a person they love — the organizer — creates a digital memorial by collecting memories from everyone who knew that person. Multiple people contribute. AI synthesizes everything. The result is a beautiful, interactive portrait of who that person was.

One-time purchase. Private. Invite only.

The three types of users

Organizer — the person who lost someone. They create the account, set up the memorial profile (name, dates, photo of the deceased), generate an invite link, share it with people who knew the person, wait for contributions to come in, then hit Generate to run the AI. They are also a contributor themselves — they go through the same flow as everyone else.

Contributor — anyone the organizer invites. They open the invite link, no account needed, select their relationship to the deceased, answer a questionnaire, upload photos from their phone, optionally upload voice recordings that contain the deceased's voice, review everything and submit. That's it — they're done.

Viewer — anyone the organizer shares the finished memorial with. Read only. They can't change anything.

The contribution flow — five steps every contributor goes through

1. Open the invite link → see the deceased's name and photo → enter their own name → click Begin
2. Select their relationship — family, friend, colleague, partner, sibling, parent, community, other
3. Answer the questionnaire — warm reflective questions designed to surface quirks, shared moments, how the person made them feel, achievements, hardships — without asking for those things directly. Can type or speak answers. Auto-saves as they go.

4. Upload photos — bulk upload from their camera roll via a mobile-optimized link. Photo metadata (timestamp, location) is preserved.
5. Upload voice recordings — audio files that contain the deceased's voice. Voicemails, voice notes, anything. Optional. Each recording gets a title from the contributor. Then review everything and confirm.

The AI pipeline — runs after organizer clicks Generate

Three pipelines run in parallel:

Pipeline 1 — Questionnaire responses Claude Sonnet reads all contributor responses together, tagged by relationship type. Finds recurring themes that multiple contributors mentioned independently without coordinating. Each theme gets a label, a summary paragraph, direct quotes from contributors, and a prominence score based on how many people signaled it.

Pipeline 2 — Photos GPT-4o Vision analyzes every photo — scene, emotion, objects, people, setting. Then matches each photo to the themes found in Pipeline 1 by comparing visual output against questionnaire text. Photos get assigned to themes. Same photo can belong to multiple themes. Photos are ordered chronologically using metadata.

Pipeline 3 — Voice recordings Whisper transcribes every recording. A separate LLM pass analyzes each transcript organically — finds recurring phrases, emotional moments, laughter, unique expressions. Pulls one key quote per recording. Assigns an AI-generated category label. Completely self-contained — doesn't feed into the other outputs.

The four outputs — what viewers see

Story tab — a linear slideshow. Photos paired with matched quotes from contributor responses. Sequenced into a narrative arc — establishing memories early, emotional depth in the middle, warm close at the end. Viewer steps through manually at their own pace. Progress bar shows where they are.

Constellation tab — an interactive theme map. Each theme from Pipeline 1 is a node. Node size reflects how many contributors mentioned it. Lines between nodes show relationship types — solid for family, dashed for friends, dotted for colleagues. Click any node and see the summary paragraph, direct quotes, and matched photos for that theme. Filter by relationship type.

Voices tab — dedicated page for the deceased's own voice. Every recording submitted. Each card shows the contributor-given title, the AI-extracted key quote, a play button, the full transcript, and the AI category label. Nothing else on this page — no photos, no AI narration. Just the person themselves.

All Photos tab — the complete photo archive. Organized into AI-generated albums named from what was actually found — not generic categories. Photos ordered chronologically within each album. Every submitted photo lives here.

The tech stack

Frontend: React + Vite + Tailwind + Framer Motion. Hosted on Vercel free tier. Pages: login, signup, dashboard, memorial create, memorial manage, contribute flow (5 pages), memorial output, share link.

Backend: Node.js + Express + TypeScript. Prisma ORM talking to Supabase Postgres. Supabase Auth for authentication. Supabase Storage for all files. Hosted on Render or Railway for demo.

Database: Supabase Postgres. 11 tables — users, memorials, invite_links, contributors, questionnaire_responses, media_assets, voice_recordings, themes, theme_quotes, theme_media_links, ai_jobs, ai_outputs.

Storage: Supabase Storage bucket called memorial-assets. Folder structure:
memorials/{memorial_id}/cover/,
memorials/{memorial_id}/contributions/{contributor_id}/photos/,
memorials/{memorial_id}/contributions/{contributor_id}/voice/,
memorials/{memorial_id}/outputs/.

AI: GPT-4o Vision for photos, Claude Sonnet for questionnaire synthesis and theme generation, Whisper for transcription.